



ECS: A PROGRAMMING PARADIGM

Azriel Hoh





ABOUT

- **Before:** Automated distributed system management.
- **Now:** Making a Rust game.





AGENDA

1. Data Organisation
 - OO: Array of Structs
 - EC: Struct of Arrays
2. Game Scenario
3. Patterns
 - Generational Arena
 - Storages
4. Logic
 - MVC
 - Systems
 - Dispatcher
5. Summary
 - Trade-offs





DATA: OO

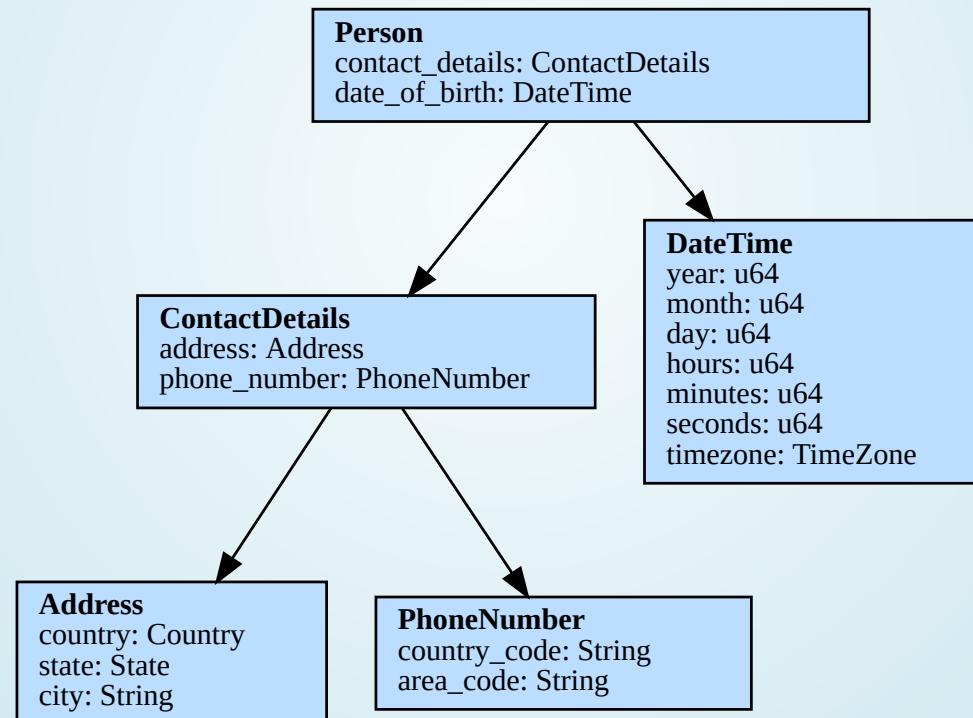




DATA: OO

We usually define data types hierarchically:

Hierarchical data organization





DATA: OO

In code, that looks something like this:

```
struct Person {  
    contact_details: ContactDetails,  
    date_of_birth: DateTime<Utc>,  
}  
  
struct ContactDetails { /* fields */ }  
struct DateTime<Tz> { /* fields */ }
```





DATA: OO

A list of people can be stored like this:

```
struct World {  
    people: Vec<Person>,  
}
```

This layout is:

- Hierarchical.
- Coined *arrays of structs* (AOS).





DATA: STRUCT OF ARRAYS





DATA: STRUCT OF ARRAYS

Split Person into parts, keep the same parts together:

```
struct World {  
    contact_details: Vec<ContactDetails>,  
    date_of_births: Vec<DateTime<Utc>>,  
}
```

Each index into the `Vec<_>`s represents a person.

Need to make sure `Vec<_>`s have the same length.

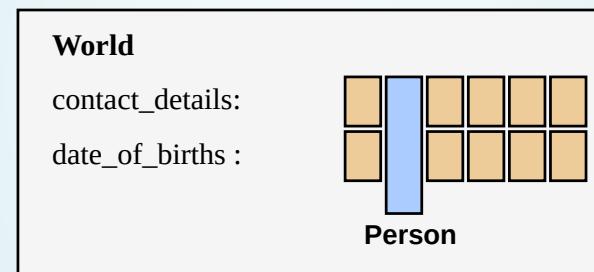




DATA: STRUCT OF ARRAYS

In picture form:

EC Data Organization (simplified)



A Person is a vertical slice from both `Vec<_>`s.

Person no longer exists, it's now an abstract concept.





DATA: STRUCT OF ARRAYS



Why would anyone want to do that?





CALCULATE AVERAGE AGE





CALCULATE AVERAGE AGE

Formula:

```
average_age = total_age / number_of_people
```





CALCULATE AVERAGE AGE

```
struct Person {
    contact_details: ContactDetails,
    date_of_birth: DateTime<Utc>,
}

let people: Vec<Person> = /* ... */;
let now = Utc::now();

let mean_age = people
    .iter()
    // Sum everyone's age
    .fold(0., |sum, p| sum + years_since(p.date_of_birth, now))
    // Divide by number of people
    / people.len() as f32
```





MEMORY ACCESS

Here is the data that we iterate over:

P0	Contact Details	Address
		Phone Number
	Date of Birth	
P1	Contact Details	Address
		Phone Number
	Date of Birth	
P2	Contact Details	Address
		Phone Number
	Date of Birth	
P3	Contact Details	Address
		Phone Number
	Date of Birth	





MEMORY ACCESS

Here is the data that we need:

P0	Contact Details	Address
		Phone Number
	Date of Birth	
P1	Contact Details	Address
		Phone Number
	Date of Birth	
P2	Contact Details	Address
		Phone Number
	Date of Birth	
P3	Contact Details	Address
		Phone Number
	Date of Birth	





CALCULATE AVERAGE AGE: SOA

```
let contact_details: Vec<ContactDetails> = /* ... */;
let date_of_births: Vec<DateTime> = /* ... */;
let world = World { contact_details, date_of_births };

let now = Utc::now();

// Note: we don't actually read `contact_details`.
let mean_age = world
    .date_of_births
    .iter()
    // Sum everyone's age
    .fold(0., |sum, dob| sum + years_since(*dob, now))
    // Divide by number of people
    / world.date_of_births.len() as f32
```





MEMORY ACCESS: SOA

Here is the data that we iterate over:

D0	Date of Birth
D1	Date of Birth
D2	Date of Birth
D3	Date of Birth





MEMORY ACCESS: SOA

Here is the data that we need:

D0	Date of Birth
D1	Date of Birth
D2	Date of Birth
D3	Date of Birth





MEMORY ACCESS: SOA

Should I care?





DATA LOCALITY





DATA LOCALITY

When you fetch data, the computer thinks:

Maybe I should fetch more!





DATA LOCALITY

	Hierarchical	SOA
Cache Capacity	P0	DateTime0
		DateTime1
	P1	DateTime2
		DateTime3
	P2	DateTime4
		DateTime5
	P3	DateTime6
		DateTime7
	P4	DateTime8
		DateTime9





BENCHMARK: AOS VS SOA



[azriel91/aos_vs_soa/lib.rs#L112-L117](https://github.com/azriel91/aos_vs_soa/lib.rs#L112-L117)

```
git clone git@github.com:azriel91/aos_vs_soa.git
cd aos_vs_soa
cargo bench
```





BENCHMARK: AOS VS SOA

Sample output:

```
running 2 tests
test tests::array_of_structs::avg_age ... bench:    2,427,309 ns/iter (+/- 201,343)
test tests::struct_of_arrays::avg_age ... bench:    2,192,988 ns/iter (+/- 136,111)

test result: ok. 0 passed; 0 failed; 0 ignored; 2 measured; 0 filtered out
```





BENCHMARK: AOS VS SOA

name	aos ns/iter	soa ns/iter	diff ns/iter	diff %	speedup
::avg_age	2,427,309	2,192,988	-234,321	-9.65%	x 1.11
::avg_age	2,438,766	2,215,783	-222,983	-9.14%	x 1.10
::avg_age	2,444,267	2,196,445	-247,822	-10.14%	x 1.11
::avg_age	2,430,973	2,195,141	-235,832	-9.70%	x 1.11
::avg_age	2,462,707	2,185,442	-277,265	-11.26%	x 1.13
::avg_age	2,414,035	2,211,378	-202,657	-8.39%	x 1.09
::avg_age	2,431,339	2,219,773	-211,566	-8.70%	x 1.10
::avg_age	2,433,808	2,225,186	-208,622	-8.57%	x 1.09
::avg_age	2,404,514	2,175,141	-229,373	-9.54%	x 1.11
::avg_age	2,409,768	2,200,761	-209,007	-8.67%	x 1.09





BREATHER





SCENARIO: GAME





SCENARIO: GAME

Disclaimer

*These examples are based on the specs crate.
There are many other ECS implementations,
whether in Rust or other languages.*





SCENARIO: GAME

Player

- Render
- Position
- Velocity
- Input

Monster

- Render
- Position
- Velocity

Speech Bubble

- Render
- Input

Map

- Render
- Bounds





OO MODEL

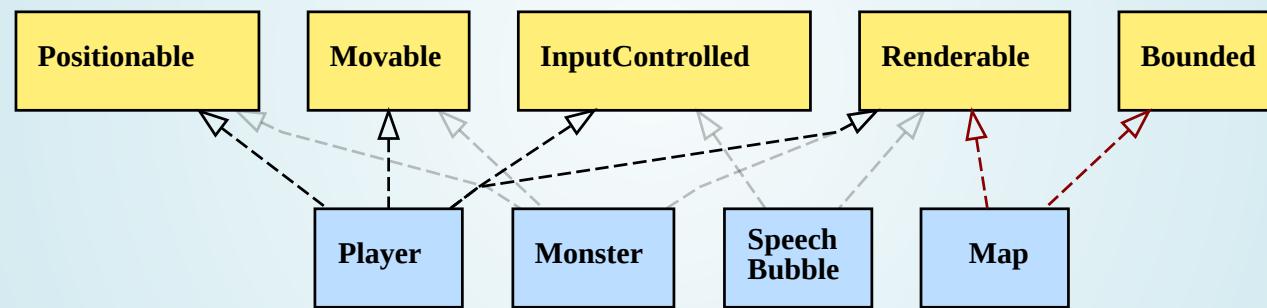
- **Good Object Oriented Design (GOOD)**
 - Interface inheritance
 - Composition
- **Object Oriented Programming Syntax (OOPS)**
 - Implementation inheritance





OO MODEL

Interface Inheritance





OO MODEL

In code:

```
trait Renderable {}
trait Positionable {}
trait Movable {}
trait InputControlled {}

struct Player;
impl Renderable for Player {}
impl Positionable for Player {}
impl Movable for Player {}
impl InputControlled for Player {}

struct Monster;
impl Renderable for Monster {}
impl Positionable for Monster {}
impl Movable for Monster {}

struct World {
    players: Vec<Player>,
    monsters: Vec<Monster>,
}
```





OO MODEL

Problems:

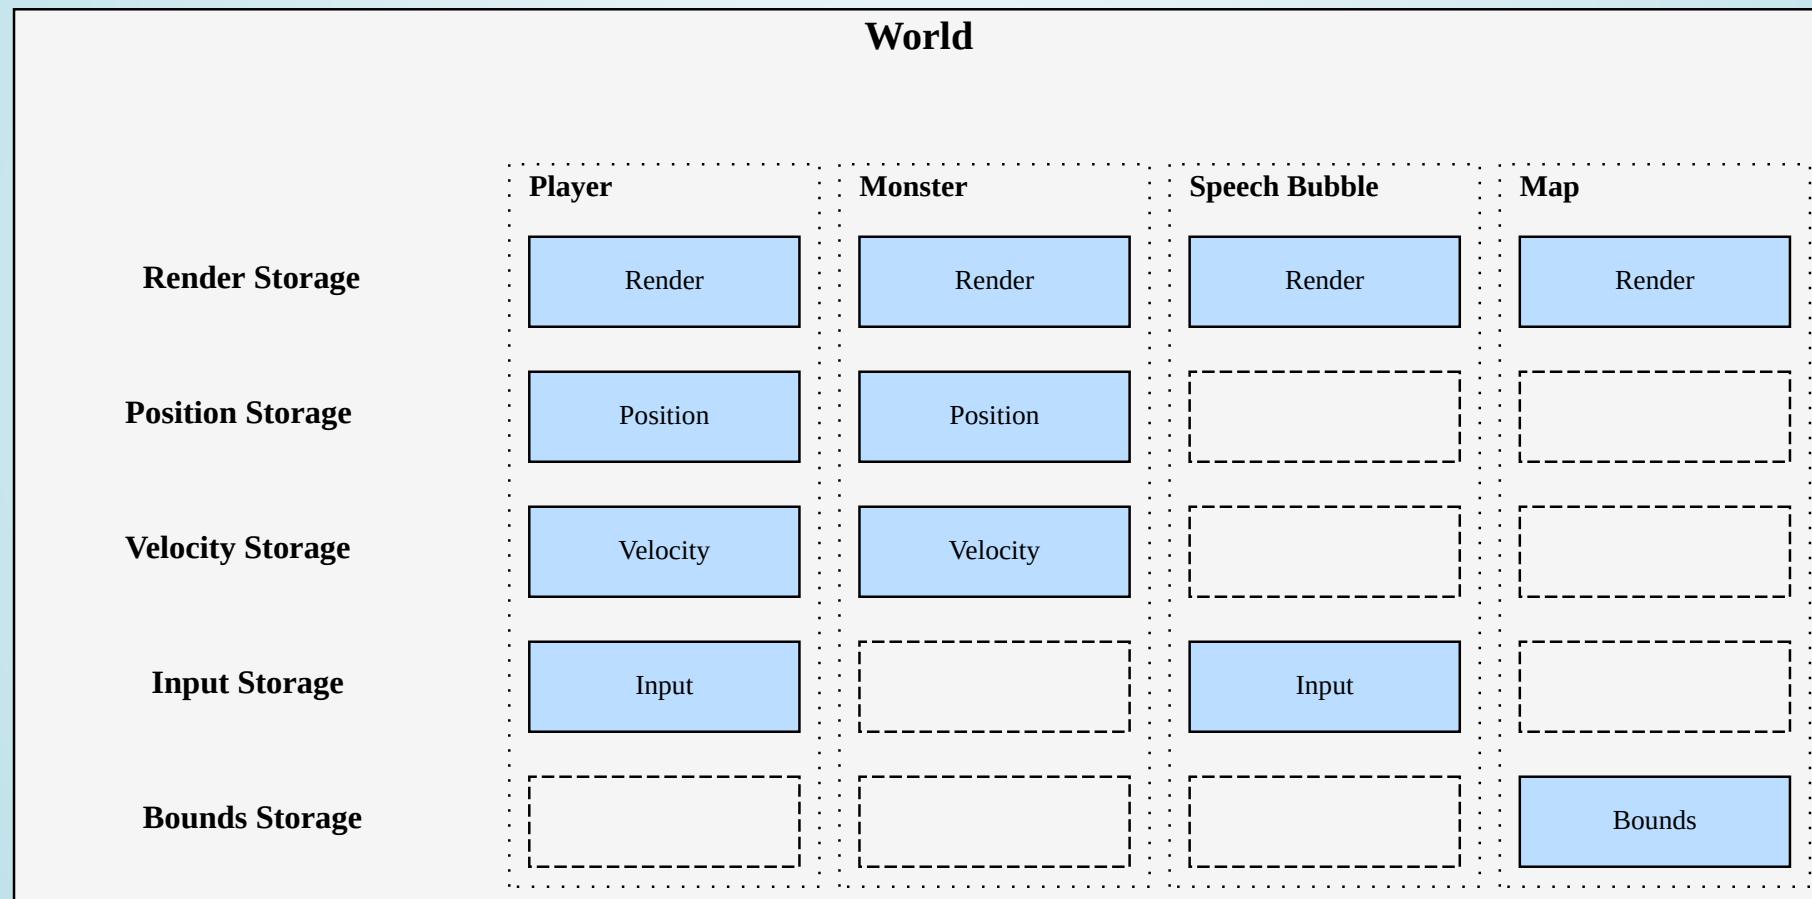
- Development overhead for common logic.
- Borrowing parts of an object.
 - Arc<Mutex<_>>
 - Copy / Clone
 - Deconstruct / Reconstruct
 - "View" structs





EC MODEL

Some Components only apply to some entities.





EC MODEL

In code, components can be represented like this:

```
struct World {
    // Note: Can't simply store traits like this,
    // In this example, components must be concrete types
    renderables: Vec<Option<Renderable>>,
    positions: Vec<Option<Position>>,
    velocities: Vec<Option<Velocity>>,
    inputs: Vec<Option<Input>>,
    bounds: Vec<Option<Bound>>,

    // How to track each type of entity? `usize`!
    players: Vec<usize>,
    monsters: Vec<usize>,
    speech_bubbles: Vec<usize>,
    maps: Vec<usize>,
}
```





EC MODEL

Problem?





EC MODEL

More like, problems!

- **Memory bloat:** Unused slots.
- **Entity creation:** Finding a free index.
 - **Overflow:** What if we run out of indices?
- **Entity deletion:**
 - **Mutable access:** Remove components from all storages.
 - **Can we delete:** Is anyone referencing this entity?





BREATHER





GENERATIONAL ARENA





GENERATIONAL ARENA

Arena? The Colosseum!





GENERATIONAL ARENA

Arena: Request a chunk of memory, and manage it yourself.

Generational: Use a *generational index* to track free space in the arena.





GENERATIONAL ARENA

Instead of referencing an entity by `usize`, use:

```
#[derive(Eq, PartialEq, ...)]
pub struct GenerationalIndex {
    index: usize,
    generation: u64,
}
```

If the generation does not match, it's a different entity.





GENERATIONAL ARENA

Key concepts:

- Keep a pool of memory, remember which slots are "empty"
- Track a generation number used when inserting data.
- When freeing a slot, ensure the generation is at least one more than the generation of the freed slot.





GENERATIONAL ARENA

Instead of:

- `Vec<T>` where `T` is the stored component
- An incrementing `usize` to track next free index

Use `Vec<Entry<T>>`, where `Entry` is:

```
enum Entry<T> {
    Free { next_free: Option<usize> },
    Occupied { generation: u64, value: T },
}
```





GENERATIONAL ARENA

Generational Arena

```
items: Occupied {  
    generation: 0,  
    data: T,  
}  
          Free {  
    next: Some(2)  
}  
          Free {  
    next: Some(3)  
}  
          Free {  
    next: Some(4)  
}  
          Free {  
    next: None  
}
```

```
// Vec index to insert element  
free_list_head: Some(1)
```

```
// Generation to insert element  
generation: 0
```

```
// Length of vec  
len: 5
```





GENERATIONAL ARENA

Generational Arena

```
items: [ Occupied { generation: 0, data: T, }, Free { next: None } ]
```

```
// Vec index to insert element  
free_list_head: Some(4)
```

```
// Generation to insert element  
generation: 0
```

```
// Length of vec  
len: 5
```





GENERATIONAL ARENA

Generational Arena

items:	Occupied { generation: 0, data: T, } Free { next: Some(4), }	Occupied { generation: 0, data: T, } Occupied { generation: 0, data: T, }	Occupied { generation: 0, data: T, } Free { next: None }
--------	--	--	--

```
// Vec index to insert element  
free_list_head: Some(1)
```

```
// Generation to insert element  
generation: 1
```

```
// Length of vec  
len: 5
```





GENERATIONAL ARENA

Solves:

- **Memory bloat:** Unused slots.
- **Entity creation:** Finding a free index.
 - **Overflow:** Reuse existing indicies
- **Entity deletion:**
 - **Mutable access:** Old gen signals absence.
 - **Can we delete:** Old gen signals absence.





EC STORAGES

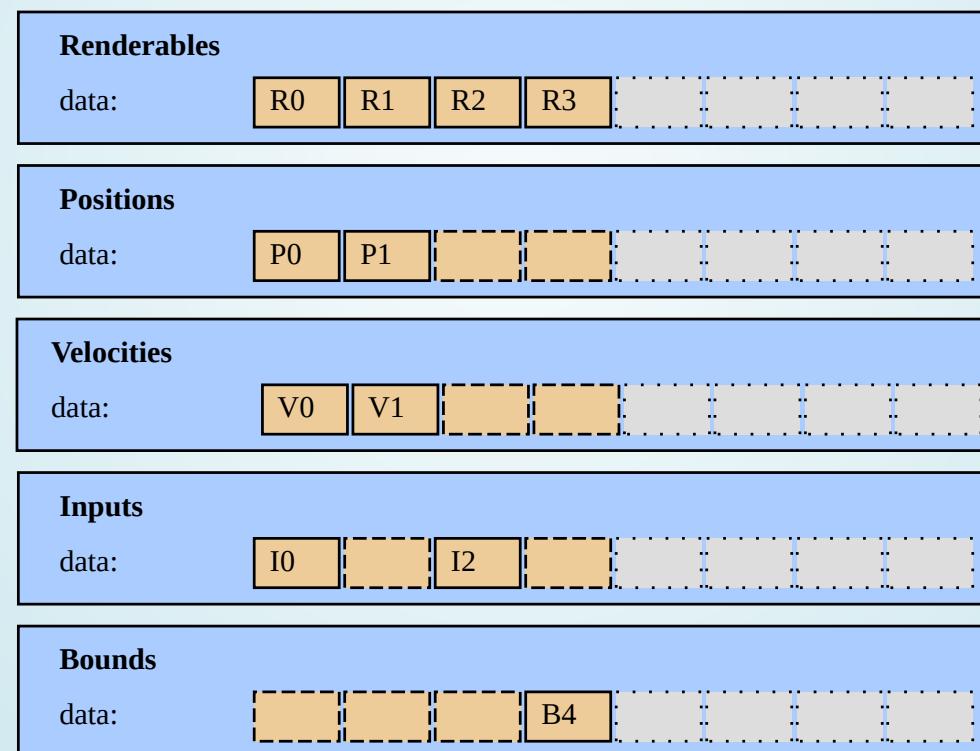




EC STORAGES

VecStorage can be a waste of memory:

VecStorage

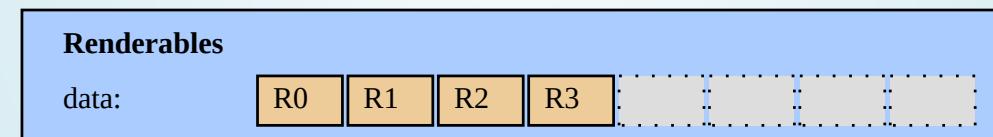




EC STORAGES

Vec is optimal for frequent components:

VecStorage

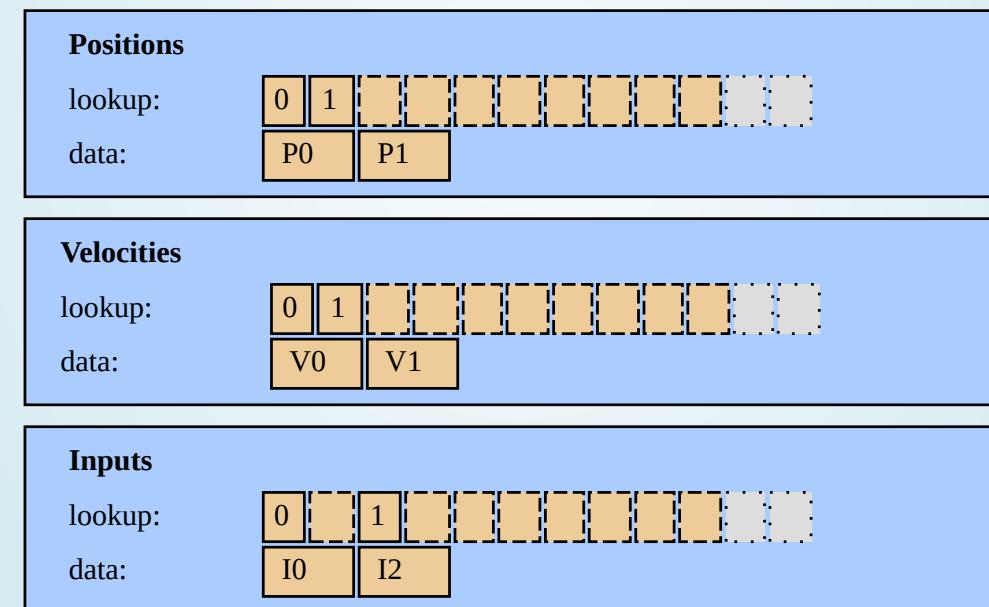




EC STORAGES

A lookup Vec saves memory for common components:

DenseVecStorage

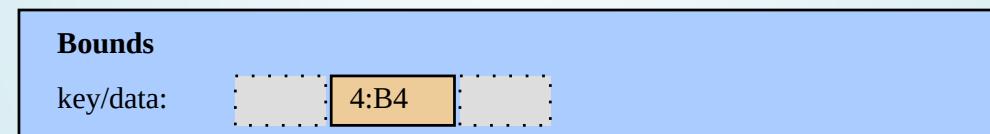




EC STORAGES

Rare components can be stored using a HashMap:

HashMapStorage





EC STORAGES

Solves:

- **Memory bloat:** Unused slots.
- **Entity creation:** Finding a free index.
 - **Overflow:** Reuse existing indicies
- **Entity deletion:**
 - **Mutable access:** Old gen signals absence.
 - **Can we delete:** Old gen signals absence.





BREATHER





LOGIC





TASK: UPDATE POSITIONS

Players and Monsters have a position and velocity.

Write a function to update position based on velocity.

```
position += velocity
```





LOGIC: OO

*Write the update function inside the class.
Hiding the code means better encapsulation.*

- every CS course (how we heard it)





LOGIC: OOPS

How to increase software maintenance costs:

```
impl PositionUpdate for Player {
    fn update(&mut self) {
        self.pos += self.vel;
    }
}

impl PositionUpdate for Monster {
    fn update(&mut self) {
        self.pos += self.vel;
    }
}
```





LOGIC: GOOD (MVC)

```
type Kinematic = i32;
impl Positionable for Player { fn pos_mut(&mut self) -> &mut Kinematic { &mut self.vel } }
impl Positionable for Monster { fn pos_mut(&mut self) -> &mut Kinematic { &mut self.vel } }
impl Movable for Player { fn vel(&self) -> &Kinematic { &self.vel } }
impl Movable for Monster { fn vel(&self) -> &Kinematic { &self.vel } }

trait GameObject: Positionable + Movable {}

fn position_update(game_objects: &mut [Box<GameObject>]) {
    for i in 0..game_objects.len() {
        let vel = game_objects[i].vel();
        *game_objects[i].pos_mut() += *vel;
    }
}
```





LOGIC: GOOD (MVC)

- Logic is generally pretty clean.
- Difficult to update objects due to borrowing rules.

Choose from:

- Arc<Mutex<_>>
- View structs - &field.
- Deconstruct, reconstruct.
- ...





SYSTEMS





SYSTEMS

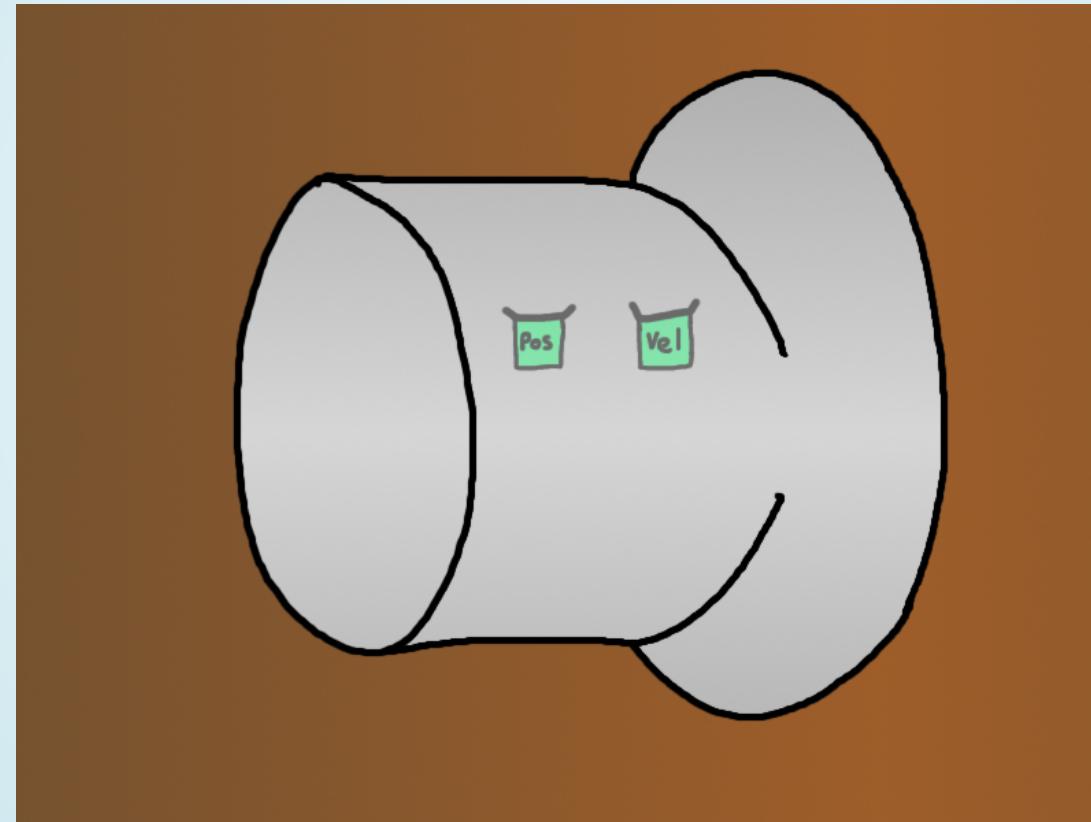
The **S** in **ECS**.





LOCK & KEY ANALOGY

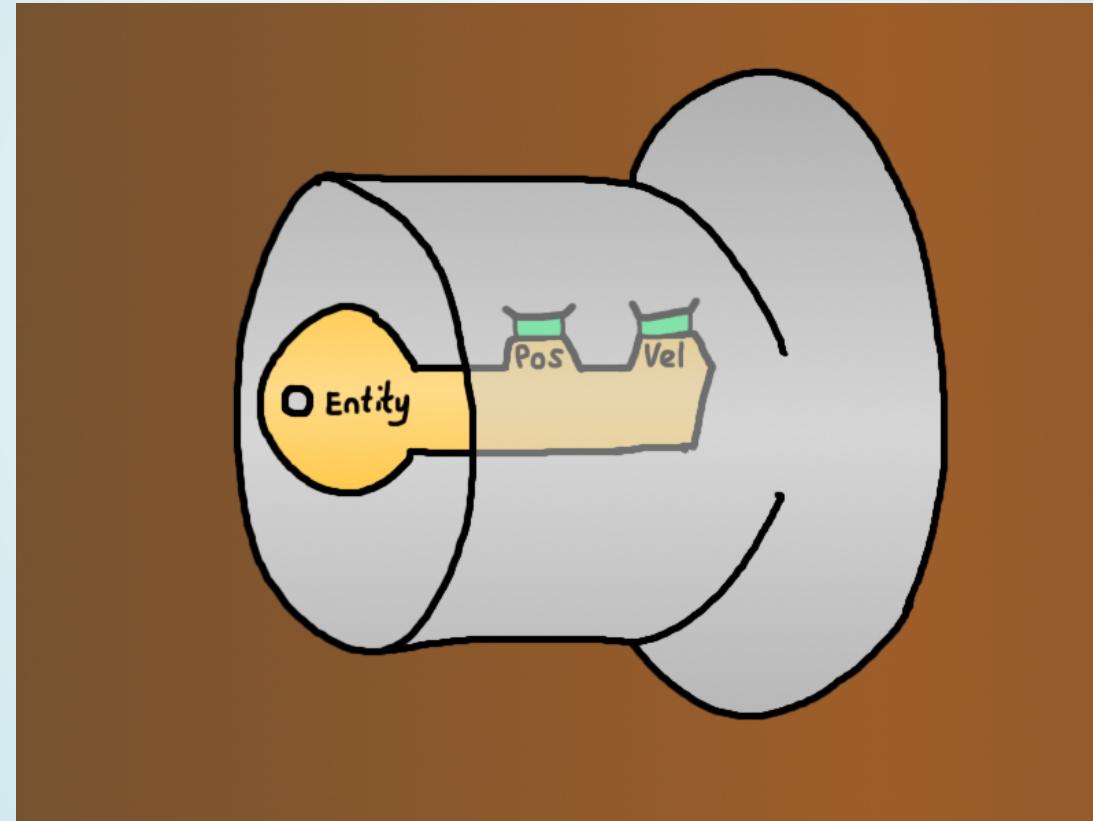
Given this lock:





LOCK & KEY ANALOGY

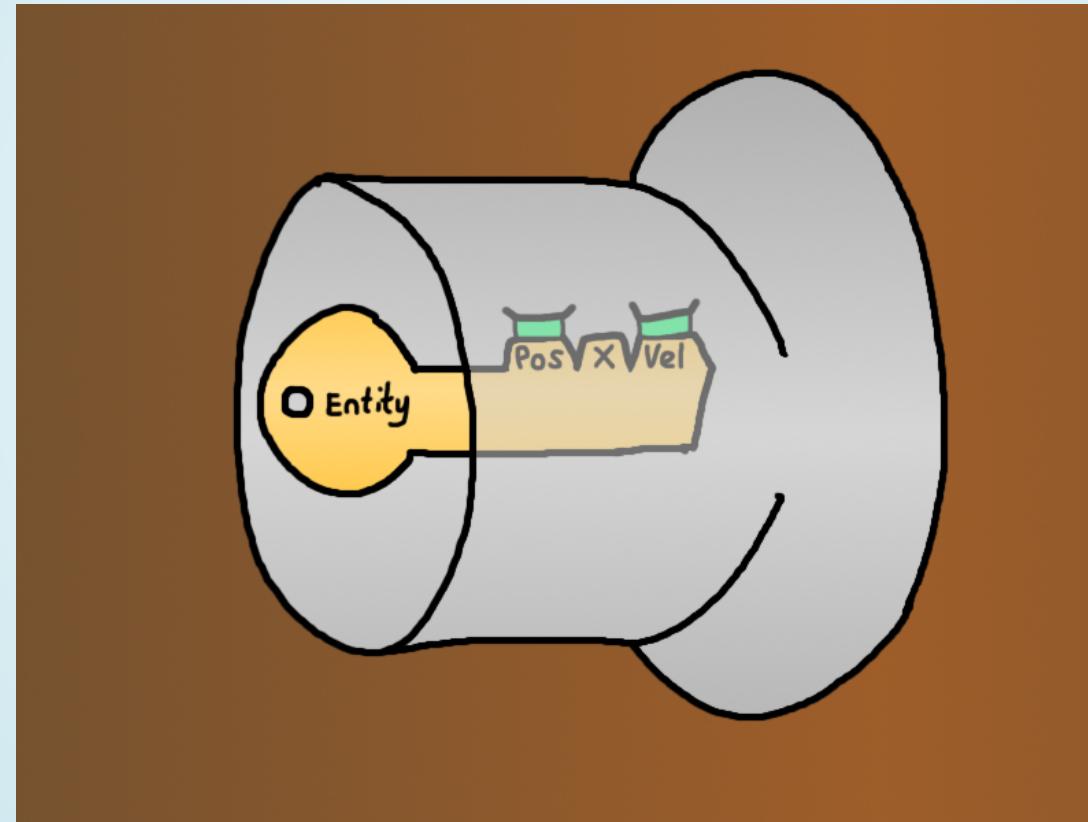
This key works:





LOCK & KEY ANALOGY

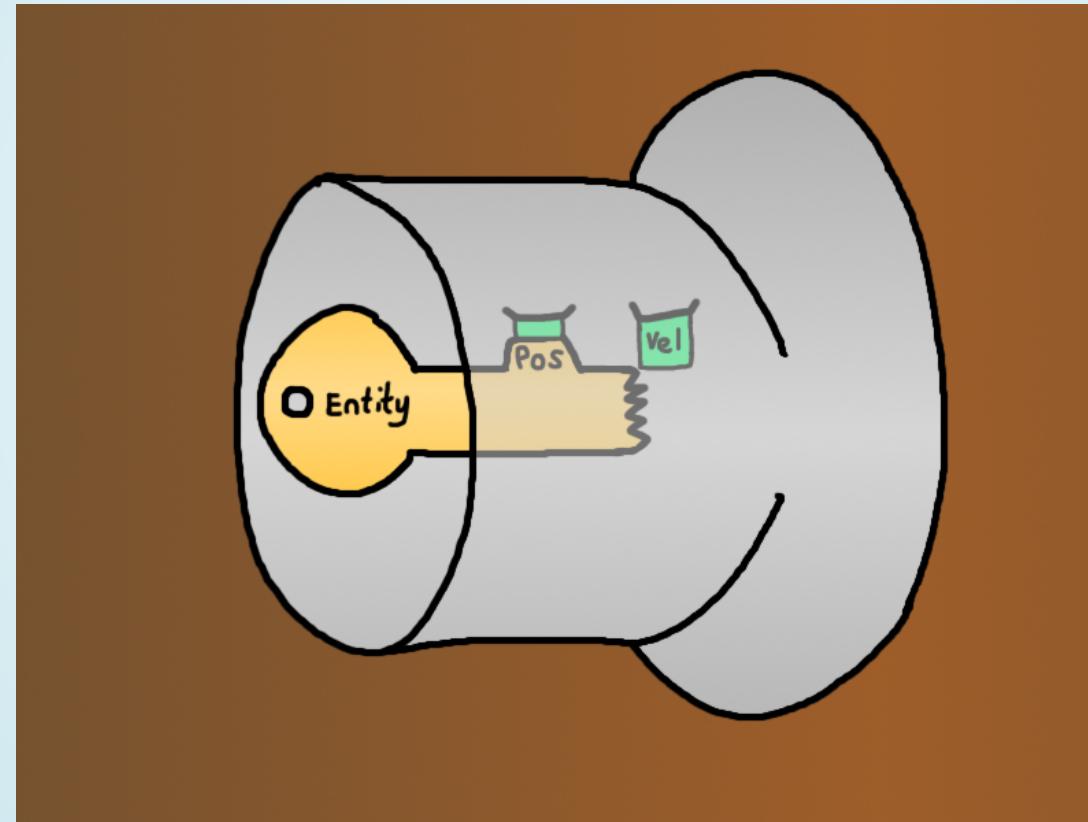
So does this:





LOCK & KEY ANALOGY

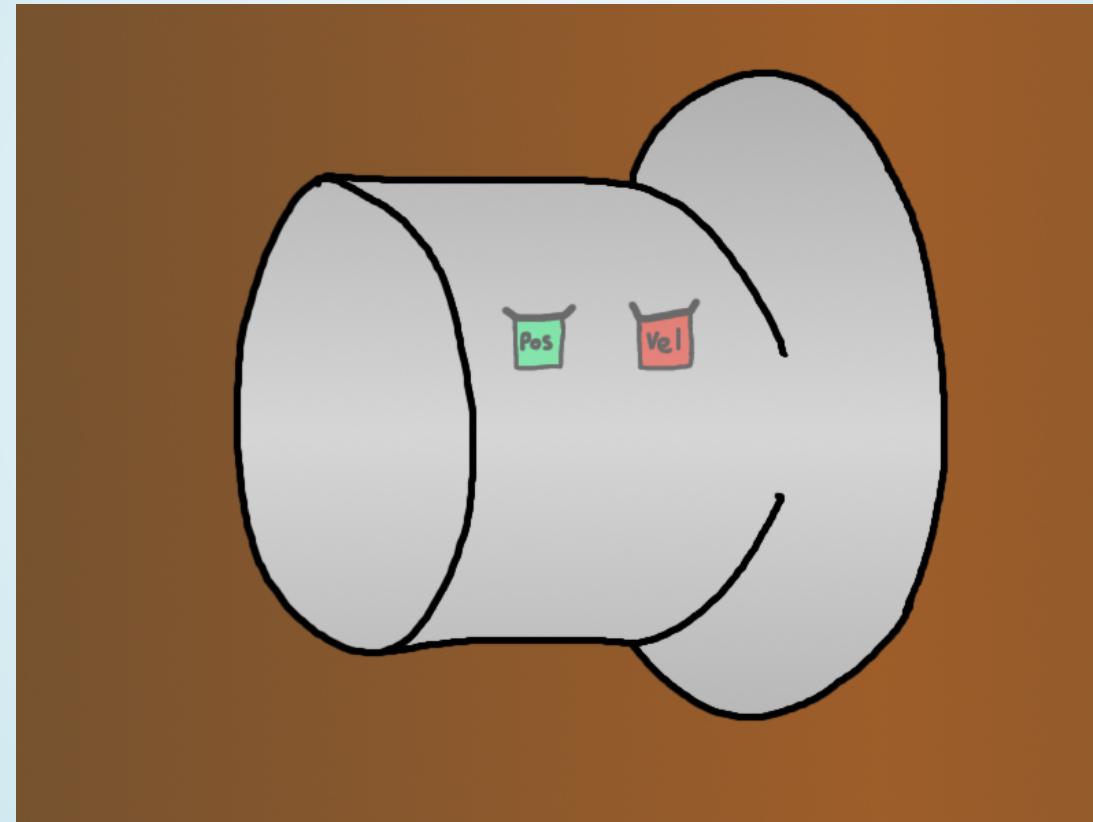
But not this:





LOCK & KEY ANALOGY

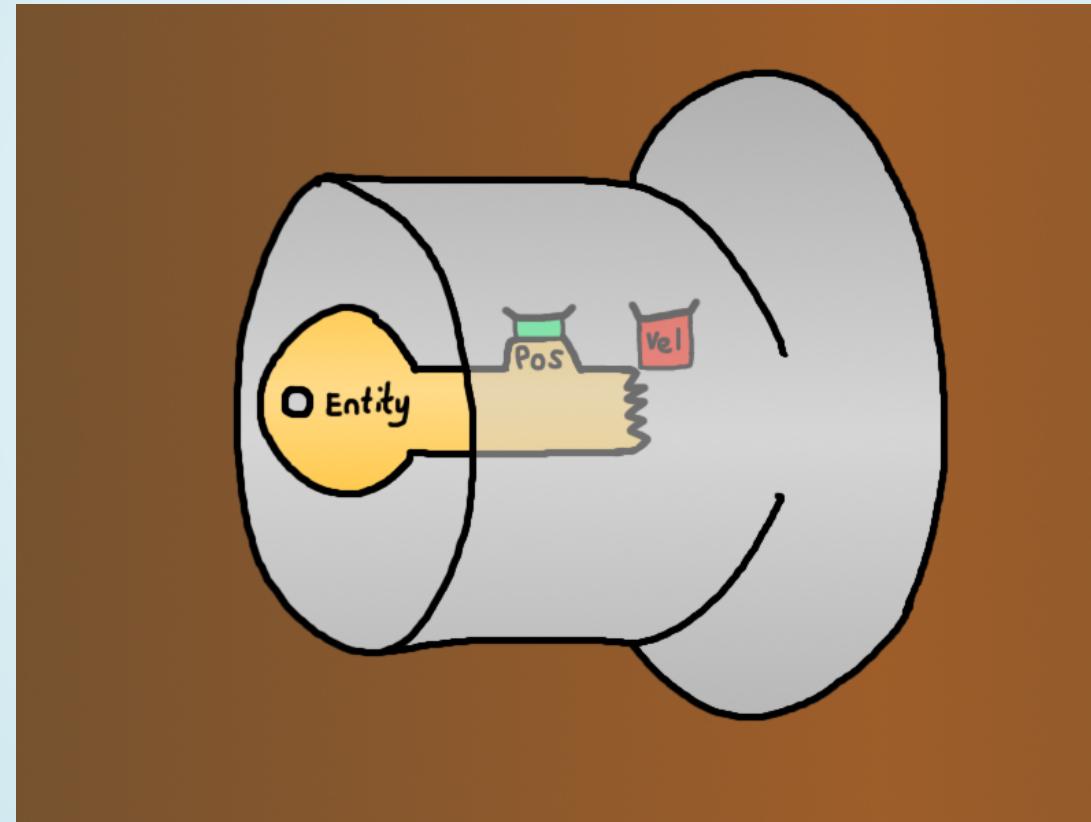
For this lock:





LOCK & KEY ANALOGY

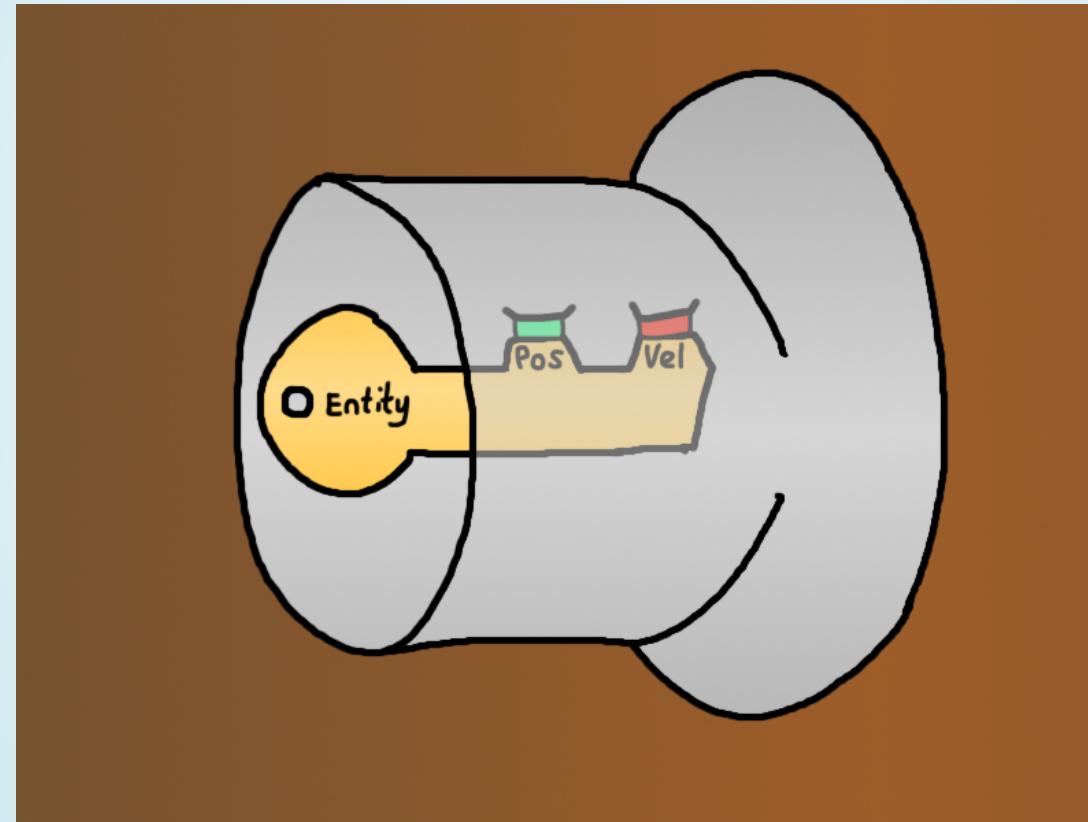
This key works:





LOCK & KEY ANALOGY

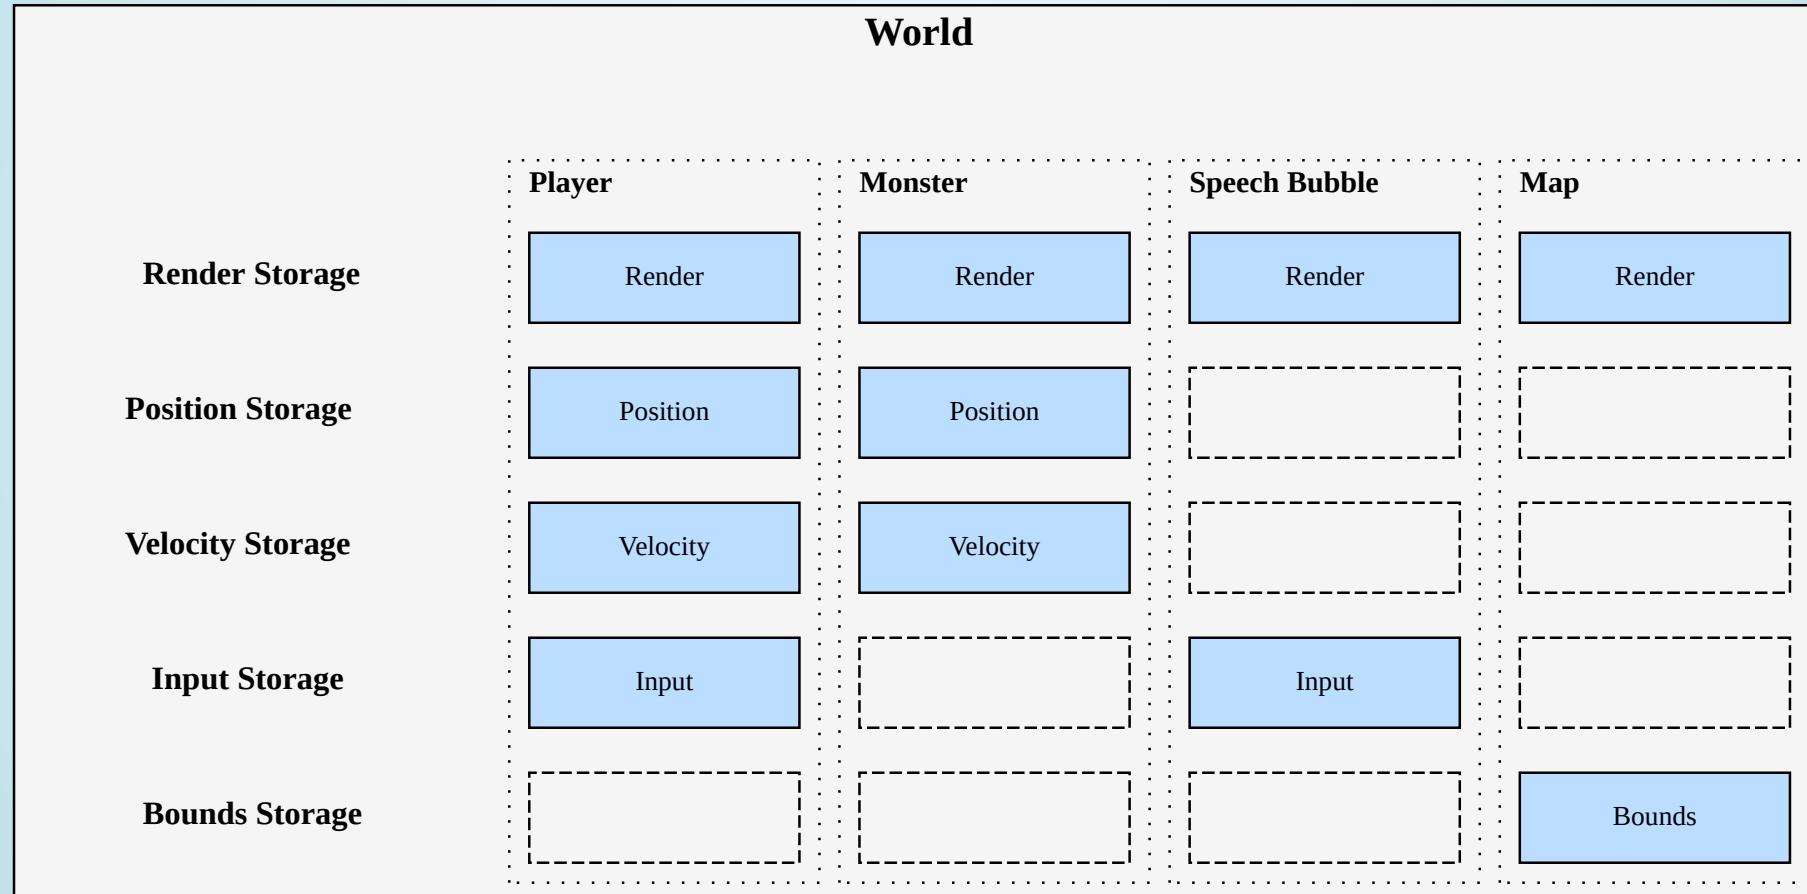
But not this:





SYSTEMS

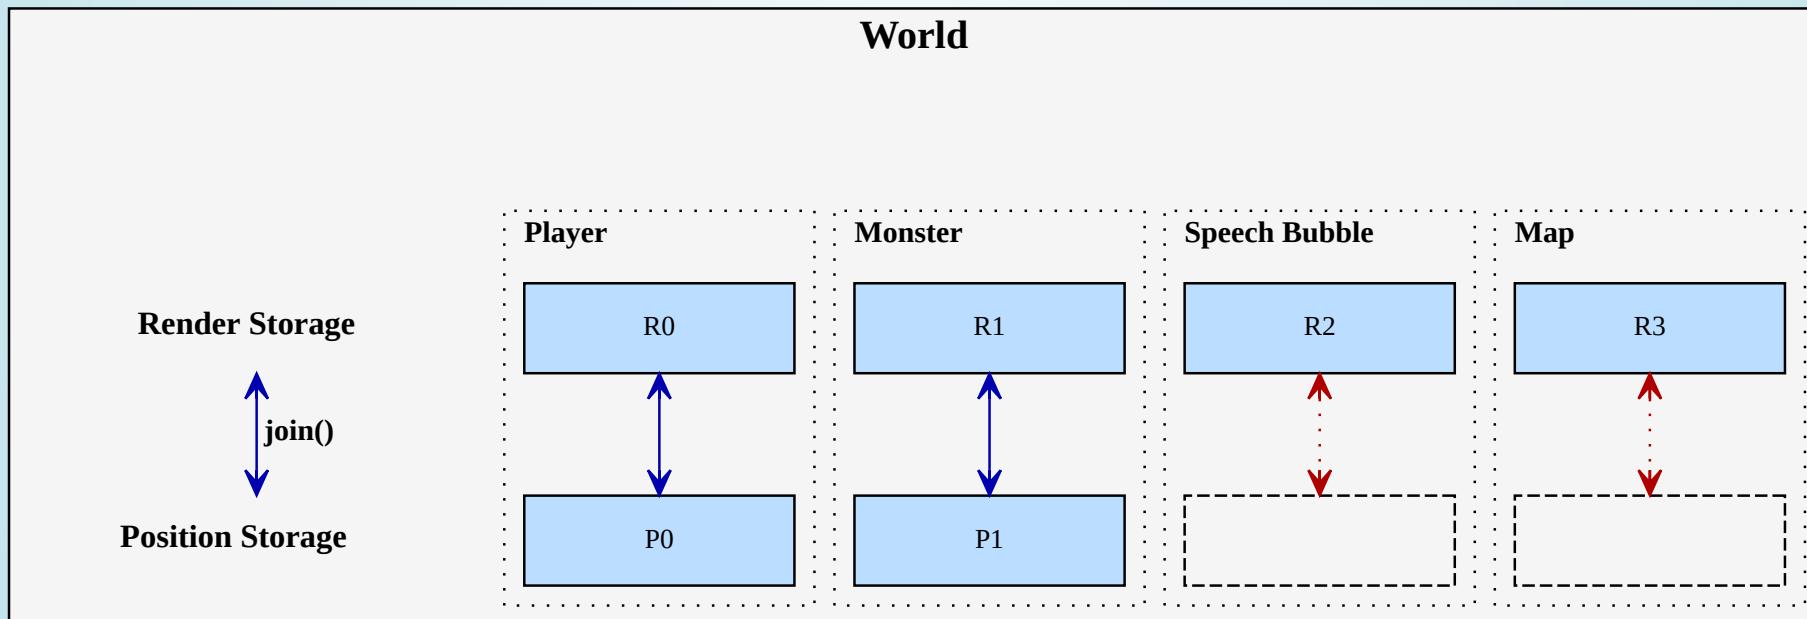
Remember this:





SYSTEMS

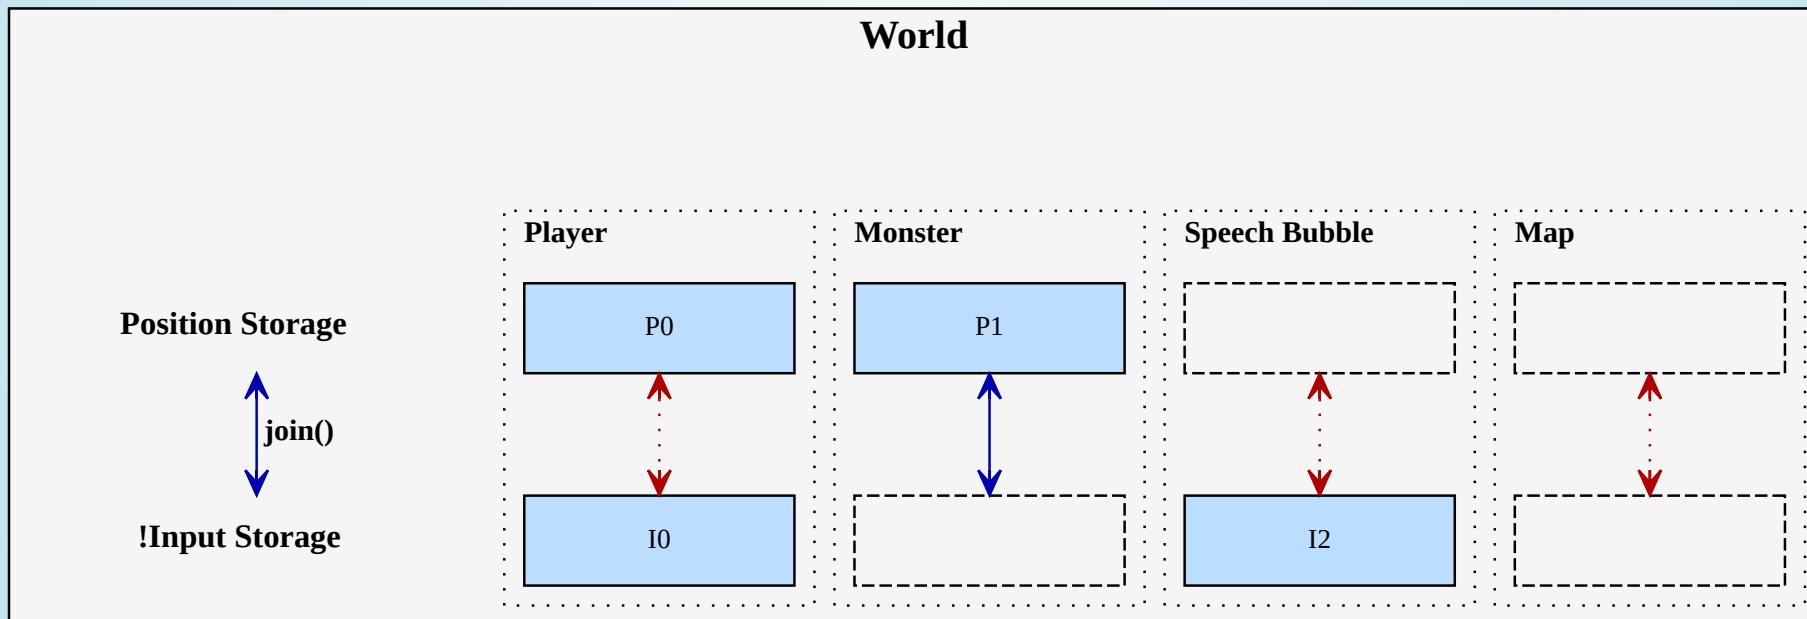
```
// Whitelist components (must haves)  
(&renderables, &positions).join() -> [(R0, P0), (R1, P1)] // player, monster
```





SYSTEMS

```
// Exclude / blacklist components  
(&positions, !&inputs).join() -> [(P1, ())] // monster
```





SYSTEMS

Compare:

```
// 00 position_update():
for i in 0..game_objects.len() {
    *game_objects[i].pos_mut() += *game_objects[i].vel();
}

// PositionUpdateSystem run():
for (pos, vel) in (&mut positions, &velocities).join() {
    *pos += *vel;
}
```





BREATHER





SYSTEM DATA





SYSTEM DATA

Systems operate over data.
We call that SystemData.

SystemData looks like this:

```
/// Systems operate on resources that are borrowed from the `World`.  
/// This is the lifetime that those references live for.  
///  
/// => The resources must live for at least as long as `'s`.  
type PositionUpdateSystemData<'s> = (  
    WriteStorage<'s, Position>,  
    ReadStorage<'s, Velocity>,  
)
```





SYSTEM DATA

And is used like this:

```
struct PositionUpdateSystem;  
  
type PositionUpdateSystemData<'s> = (  
    WriteStorage<'s, Position>,  
    ReadStorage<'s, Velocity>,  
);  
  
impl System<'s> for PositionUpdateSystem {  
    type SystemData = PositionUpdateSystemData<'s>;  
  
    fn run(&mut self, (mut positions, velocities): Self::SystemData) {  
        // Here is the system logic!  
        for (pos, vel) in (&mut positions, &velocities).join() {  
            *pos += *vel;  
        }  
    }  
}
```





DISPATCHER





DISPATCHER

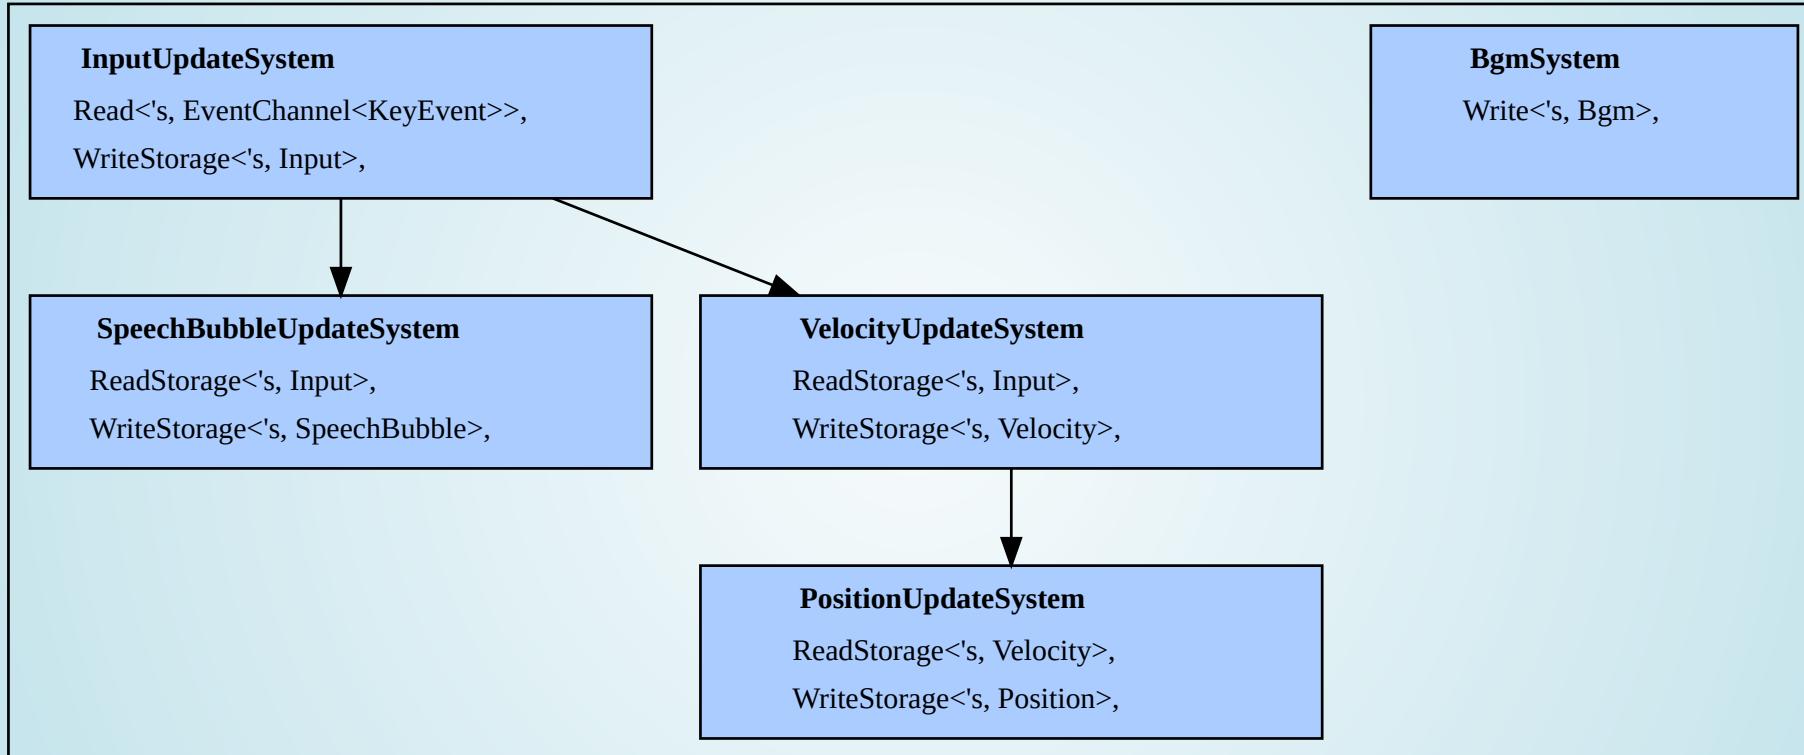
What it is:

- System graph
- Thread pool
- Execution
 - Parallel
 - Safe: multiple readers / exclusive writer





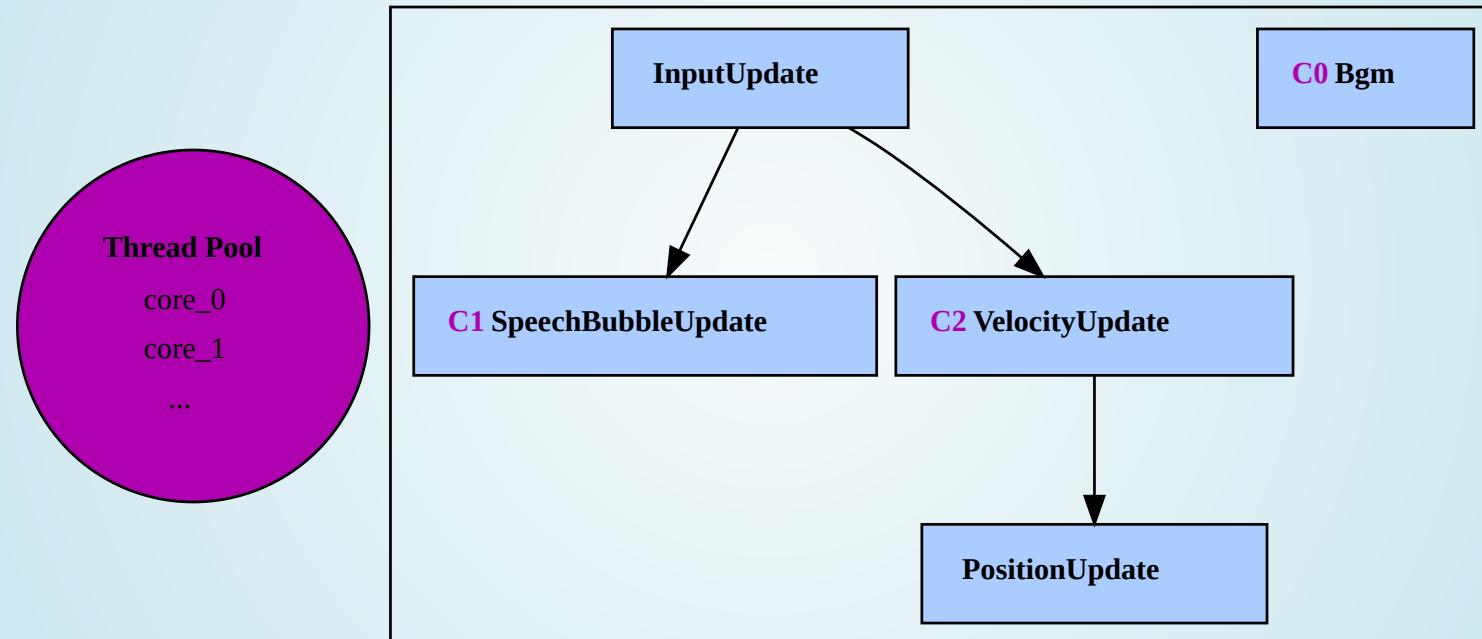
DISPATCHER: SYSTEM GRAPH



[specs/examples/full.rs#L221-L230](#)



DISPATCHER: THREAD POOL





SUMMARY

- Compare GOOD with ECS, not OOPS with ECS
- It's a trade-off.

	OO	EC
Logical partitioning	State	Behaviour
Visible concrete model		
Optimized for cache usage ¹		
Borrow-checker management ²		

¹ Rust doesn't auto-vectorize floats well.

² Ease of passing data for parallelization



QUESTIONS, ANSWERS AND COMMENTS





THANKS!

- Catherine West / @Kyrenite: Learnt *a lot* from her blog.
- The specs team: Who have written such a great library.
- Maik Klein: For expanding my knowledge on auto-vectorization.
- @medusacle: For much feedback on improving these slides.





LINKS

- Slides: https://github.com/azriel91/ecs_paradigm
- RustConf 2018 Closing Keynote: <https://www.youtube.com/watch?v=P9u8x13W7UE>
- @Kyrenite's Blog: <https://kyren.github.io/2018/09/14/rustconf-talk.html>
- specs: <https://crates.io/crates/specs>
- soa-derive: <https://github.com/lumol-org/soa-derive>
- Benchmark comparison: https://github.com/azriel91-aos_vs_soa
- Component Graph System: <https://github.com/kvark/froggy/wiki/Component-Graph-System>

